

# RapidMiner

## Linked Open Data Extension

Manual

Version 1.3.3, 05/15/14



Heiko Paulheim

Petar Ristoski

Evgeny Mitichkin

Christian Bizer

*University of Mannheim*

*Data and Web Science Group*

Contact: use the Google Group

<https://groups.google.com/forum/#!forum/rmlod>  
[rmlod@googlegroups.com](mailto:rmlod@googlegroups.com)

## Table of Contents

1 Introduction.....	3
2 Enabling the RapidMiner LOD Extension.....	4
3 Configuring SPARQL Endpoints.....	6
4 Using Linked Open Data as Input for Data Mining.....	9
5 Linking your Data to Existing Datasets.....	11
5.1 The Pattern-based Linker.....	13
5.2 The Label-based Linker.....	13
5.3 The DBpedia Lookup Linker.....	14
5.4 The DBpedia Spotlight Linker.....	16
5.5 The SameAs Linker.....	16
5.6 Combining Linkers.....	17
6 Using generators.....	18
6.1 Direct Types Generator.....	18
6.2 Data Properties Generator.....	19
6.3 Unqualified and Qualified Relation Generators.....	19
6.4 The Specific Relation Generator.....	19
6.5 The Custom SPARQL Generator.....	20
6.6 Combining Generators.....	21
7 Filtering Attributes.....	22
8 Storing Results.....	23
9 Migration from Earlier Versions of the Extension.....	24
10 Release Notes.....	25

If you use the RapidMiner Linked Open Data Extension for scientific purposes, please kindly include a reference to the following paper, which describes the underlying algorithms:

H. Paulheim, J. Fürnkranz (2012): *Unsupervised Generation of Data Mining Features from Linked Open Data*. In: International Conference on Web Intelligence, Mining, and Semantics (WIMS'12).

## 1 Introduction

Linked Open Data<sup>1</sup> is a collection of freely accessible, machine interpretable data on the web. It uses semantic web standards like RDF for data representation. As of today, there are more than 200 datasets from various domains, encompassing general knowledge as well as specific domains such as government, geographic, and life science data, scientific publications, and media. A detailed list of datasets that are available as Linked Open Data can be obtained at <http://lod-cloud.net/>.

Essentially, Linked Open Data contains data as graphs with labeled edges, where each entity is represented by a dereferencable URI. Each statement can be understood as a triple with subject, predicate, and object, such as

<<http://dbpedia.org/resource/Germany>> <<http://dbpedia.org/ontology/capital>>  
<<http://dbpedia.org/resource/Berlin>> .

The RapidMiner Linked Open Data extension leverages those data sources both to create useful datasets for Data Mining, as well as for adding background knowledge to existing datasets. For example, on a dataset for countries, data like the population, area, and GDP can be added from DBpedia<sup>2</sup> in order to improve the results of the data mining process.

In particular, the extension works with datasets that provide a SPARQL endpoint, i.e., a web service which delivers data using the query language SPARQL. A list of such datasets is available at [datahub.io](http://datahub.io).

Furthermore, data from Linked Open Data sources can be used as input to the data mining process, e.g., reading data from the Eurostat Linked Open Data endpoint and running analysis on the data with RapidMiner.

**Important notice: The RapidMiner LOD Extension usually accesses live data from the web. This can result in longer operator runtimes.**

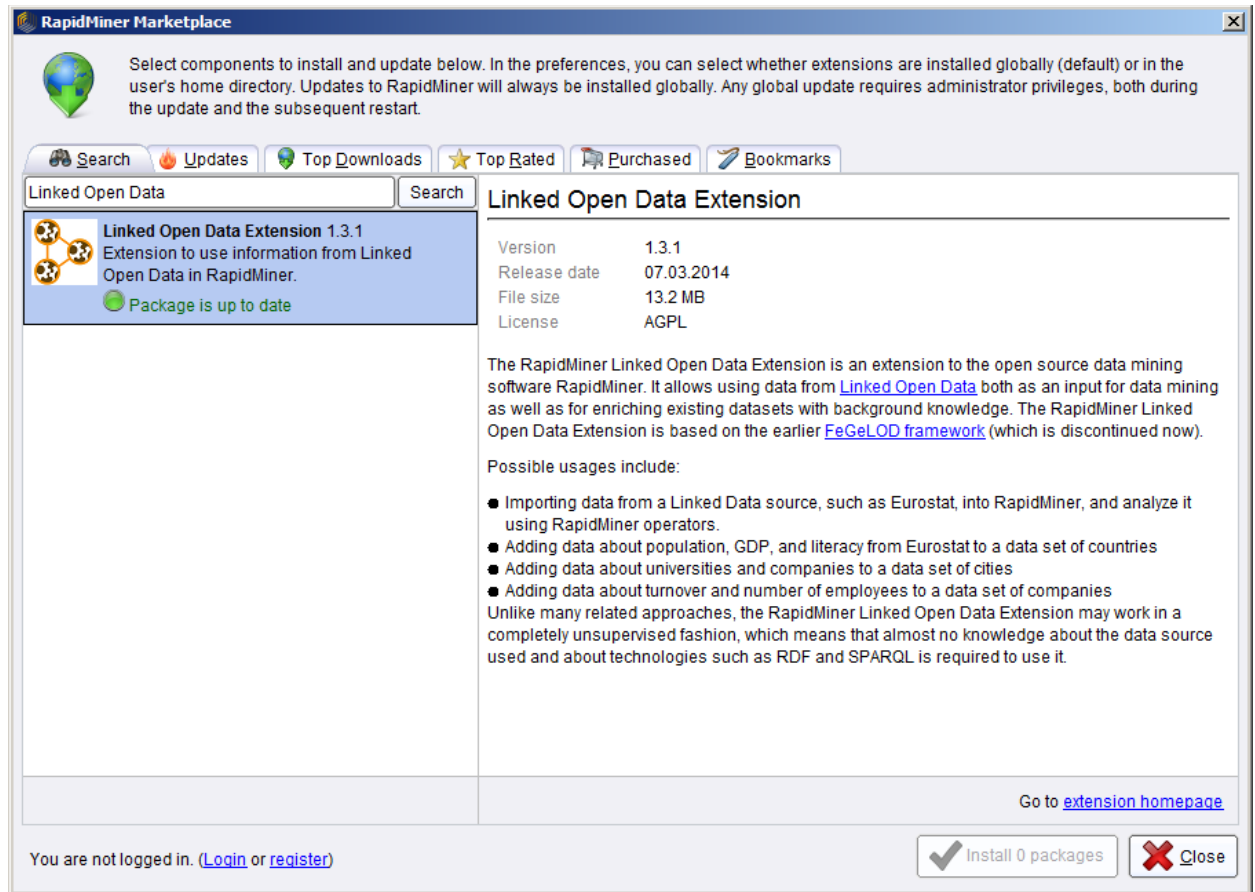
---

<sup>1</sup><http://lod-cloud.net/>

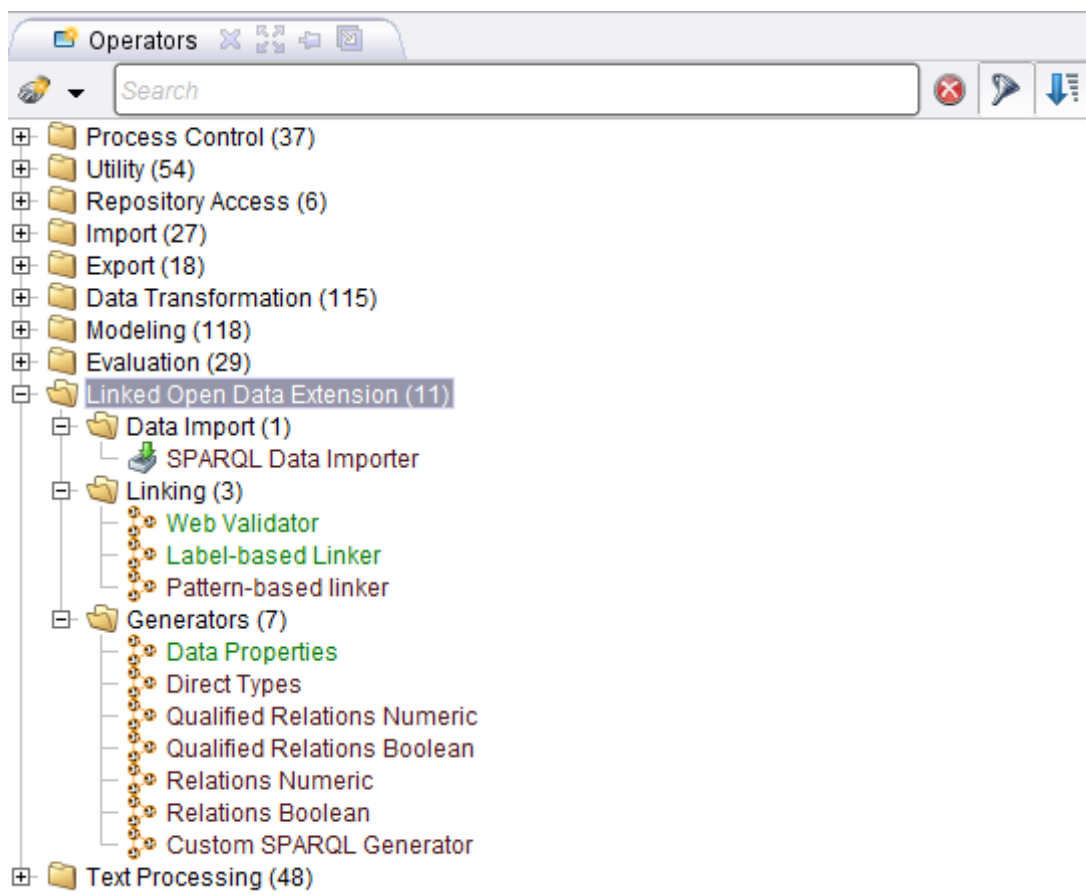
<sup>2</sup><http://dbpedia.org/>

## 2 Enabling the RapidMiner LOD Extension

1. Searching for “Linked Open Data” in the RapidMiner Marketplace dialog will find the extension and allow you to install it:



2. Create a new process. You will see a new entry (Linked Open Data Extension) in the catalog available:



### 3 Configuring SPARQL Endpoints

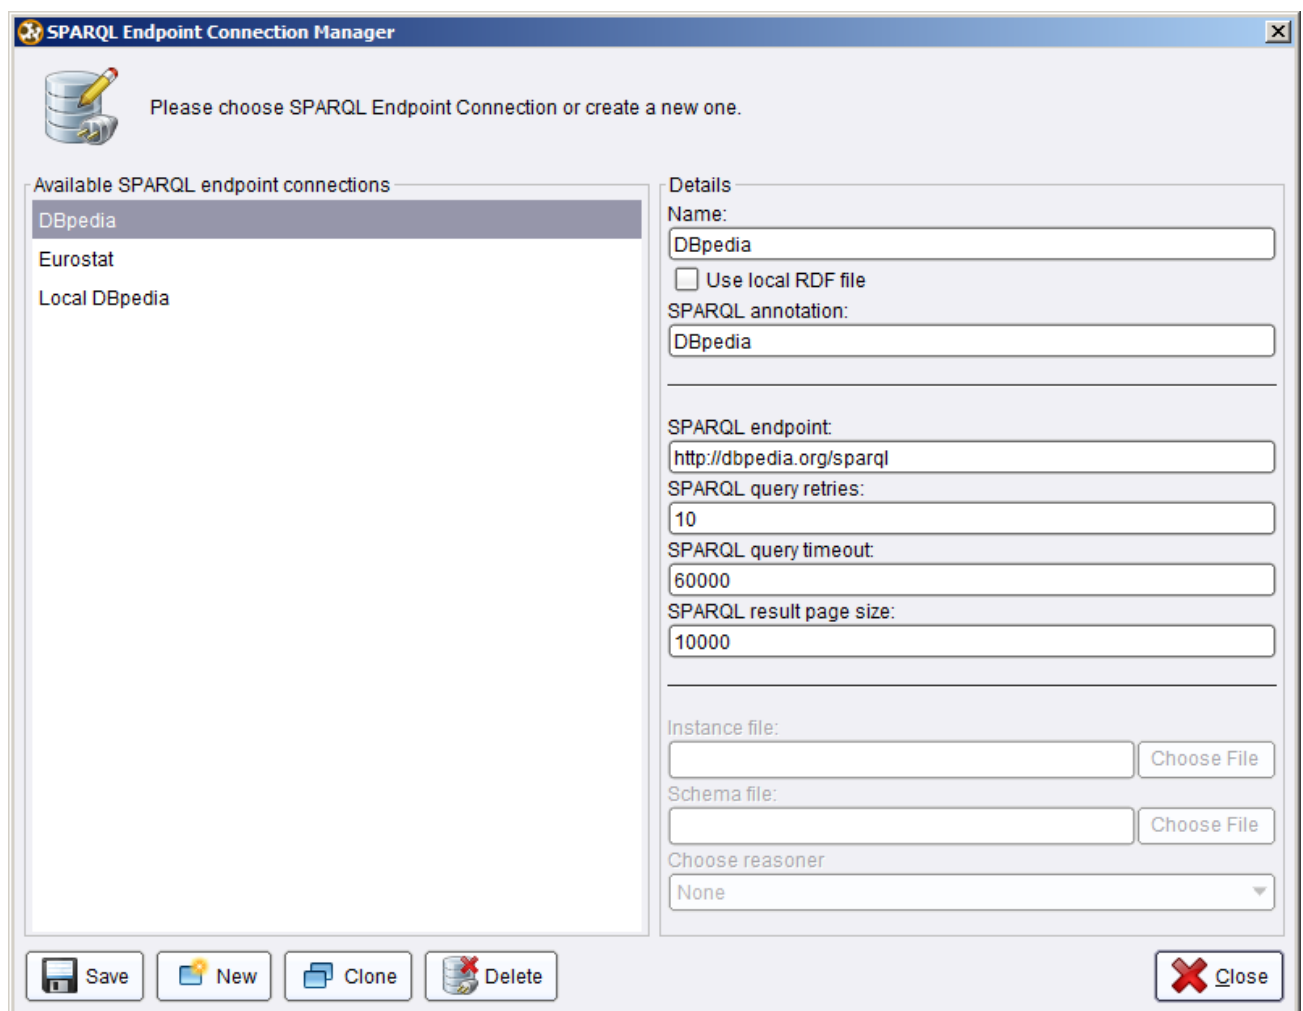
Most Linked Open Data sources provide a SPARQL endpoint. These endpoints allow for querying the underlying dataset using a standardized query language, i.e., SPARQL<sup>3</sup>.

Since most of the operators of the Linked Open Data extension use SPARQL, you will first have to define SPARQL Endpoints for the datasets you are planning to use.

To manage SPARQL Endpoints, go to the menu “Tools”->“Manage SPARQL Endpoint Connections...”. The SPARQL Endpoint Configuration Dialog Opens:

You can create new SPARQL Endpoints, as well as change the settings for existing ones.

The configuration parameters for SPARQL Endpoints are:



The screenshot shows the "SPARQL Endpoint Connection Manager" dialog box. It has a title bar with a gear icon and the text "SPARQL Endpoint Connection Manager". Below the title bar is a message: "Please choose SPARQL Endpoint Connection or create a new one." with an icon of a database cylinder and a pencil. The dialog is divided into two main sections. The left section, titled "Available SPARQL endpoint connections", contains a list box with three items: "DBpedia", "Eurostat", and "Local DBpedia". The right section, titled "Details", contains several input fields and a checkbox. The "Name:" field is filled with "DBpedia". There is a checkbox labeled "Use local RDF file" which is unchecked. The "SPARQL annotation:" field is filled with "DBpedia". The "SPARQL endpoint:" field is filled with "http://dbpedia.org/sparql". The "SPARQL query retries:" field is filled with "10". The "SPARQL query timeout:" field is filled with "60000". The "SPARQL result page size:" field is filled with "10000". Below these fields are three more fields: "Instance file:", "Schema file:", and "Choose reasoner:". Each of the first two has a text input field and a "Choose File" button. The "Choose reasoner:" field is a dropdown menu currently showing "None". At the bottom of the dialog are five buttons: "Save", "New", "Clone", "Delete", and "Close".

- Name: the SPARQL Endpoint will be selectable later by that name
- SPARQL Endpoint: the URL of the service

<sup>3</sup><http://www.w3.org/TR/sparql11-overview/>

- SPARQL Annotation: a string that will be used for naming attributes generated from this SPARQL endpoint

Three more expert parameters are available (i.e., they appear only if RapidMiner is executed in expert mode):

- Paging Size: many SPARQL endpoints have limitations w.r.t. to the amount of data they deliver. Paging ensures that all data is retrieved. For example, with a paging size value of 1000, a SELECT query is broken down into several queries like  

```
SELECT ... LIMIT 1000 OFFSET 0
SELECT ... LIMIT 1000 OFFSET 1000
SELECT ... LIMIT 1000 OFFSET 2000
```

 etc.  
 A Paging Size value of 0 disables paging.
- Query timeout defines the time which a SPARQL endpoint is given to respond (in milliseconds)
- Query retries defines the number of times a query is retried. If there are temporary errors, e.g., network problems, simply re-running the same query often resolves the issue.

Once the SPARQL Endpoint is defined, it may be used in subsequent operators. SPARQL Endpoint definitions for DBpedia and Eurostat are included as preconfigured endpoints.

The RapidMiner Linked Open Data extension also allows for local RDF files instead of remote SPARQL endpoints. To use a local RDF file, activate the corresponding checkbox. The options for local RDF files are:

- Specifying an instance and (optionally) a schema file
- Selecting a reasoner. The Jena built-in reasoners for RDFS as well as OWL Mini, OWL Micro and OWL Full are supported.<sup>4</sup>

There is also the possibility to use a default model for URL based access to Linked Open Data, i.e., accessing the data by dereferencing URIs:

---

<sup>4</sup>See <http://jena.apache.org/documentation/inference/>

**SPARQL Endpoint Connection Manager**

Please choose SPARQL Endpoint Connection or create a new one.

Available SPARQL endpoint connections:

- DBpedia
- Eurostat
- URLbasedModel**

Name: URLbasedModel

Set the model type: Use URL data

SPARQL annotation: URLbasedModel

☒ Endpoint supports COUNT

☐ Endpoint supports Property Paths

SPARQL endpoint:

SPARQL query retries: 0

SPARQL query timeout: 0

SPARQL result page size: 0

Instance file: Choose File

Schema file: Choose File

Choose reasoner


Save New Clone Delete Close

This is done by selecting “Use URL data”. You will only need one of those models for all possible LOD datasets accessed via derefencing URIs.



## 4 Using Linked Open Data as Input for Data Mining

The RapidMiner Linked Open Data Extension provides an import operator to read data from a Linked Open Data source into a RapidMiner table. The SPARQL Data Importer uses a defined SPARQL Endpoint and a custom query to generate a RapidMiner table:



The SPARQL query is defined as follows:



The table which is generated in the above example has three columns, named after the SPARQL query variables, i.e., country, GDP, and electricity. The output of the operator thus looks as follows:

Result Overview

ExampleSet (SPARQL Data Importer)

☒ Data View
☐ Meta Data View
☐ Plot View
☐ Advanced Charts
☐ Annotations

ExampleSet (24 examples, 0 special attributes, 3 regular attributes)

Row No.	country	GDP	electricity
1	Euro Area ( Ea-13 )	7775004.900	1667756.600
2	Belgique- Belgie	289508.500	79166.200
3	Ceska Republika	87205.200	52196
4	Danmark	196158.400	32552
5	Deutschland	2207200	528332
6	Eesti	9375.400	6526
7	Ireland	147569.200	15980
8	Ellada	168417.200	42273
9	España	840106	163439
10	France	1659020	350410
11	Italia	1388870.400	260809
12	Latvija	11156.600	4745
13	Lietuva	18125.800	7275
14	Luxembourg (grand-duché	26996.100	4895
15	Magyarország	82302.700	30082
16	Nederland	489854	85610
17	Österreich	235818.600	47905
18	Polska	203951.600	127160
19	Portugal	143477.900	28851
20	Roumanie	26222.200	10664

In the meta data view, you will observe that the operator automatically assigns meaningful data types to the values read from the SPARQL Endpoint. In that case, the country name is a text attribute, while GDP and electricity are numeric:

<input type="radio"/> Data View	<input checked="" type="radio"/> Meta Data View	<input type="radio"/> Plot View	<input type="radio"/> Advanced Charts	<input type="radio"/> Annotations	
ExampleSet (24 examples, 0 special attributes, 3 regular attributes)					
Role	Name	Type	Statistics	Range	Missings
regular	country	text	mode = Euro Area ( Ea-13 ) (1),	Euro Area ( Ea-13 ) (1), Belgique	0
regular	GDP	numeric	avg = 758619.800 +/- 1622566.	[9375.400 ; 7775004.900]	0
regular	electricity	numeric	avg = 170045.833 +/- 345200.2	[4745.000 ; 1667756.600]	0

## 5 Linking your Data to Existing Datasets

For using Linked Open Data as background knowledge for existing datasets, the first step is to link your data to a dataset from which you want to use additional data. For example, your dataset may contain a column “Country” with values “Germany”, “France”, etc. Linking your dataset to Linked Open Data means that an additional column is added, which contains URIs identifying those entities in Linked Open Data, such as <http://dbpedia.org/resource/Germany>, <http://dbpedia.org/resource/France>, etc.

The current version of the RapidMiner Linked Open Data extension comprises two linking methods:

1. The Pattern-based Linker builds links using a defined URI pattern. For example, the contents of the column “Country” are appended to a fixed string, e.g., “<http://dbpedia.org/resource/>”.
2. The Label-based linker searches for the contents of a column (and optionally word n-grams) in a data source, and finds the most suitable resource based on their labels. For example, for a string “United States of America”, the linker would search for word n-grams such as “United States” as well, and from all the results, use the best suited one.


Each linker has two output ports:

1. App contains the data with the new column depicting the URI
2. Att contains the attribute name which holds the URIs

The column is named ORIGINAL\_COLUMN\_link\_to\_SPARQL\_ENDPOINT\_ANNOTATION, where ORIGINAL\_COLUMN is the name of the column where the values come from (e.g., City), and SPARQL\_ENDPOINT\_ANNOTATION is the value set when configuring the SPARQL endpoint. For example, linking a column City to DBpedia will result in a new column City\_link\_to\_DBpedia.

The Att port has to be connected to subsequent steps, e.g., generators, to inform them on which attribute to perform data generation.

Since some naïve linkers (such as the Label-based Linker) may generate links that actually do not exist, the Web Validator checks the existence of each link and removes instances for which the link does not exist. Furthermore, it removes all instances for which no link could be found. The full pipeline of linking looks as follows:



The Web Validator has three outputs: App is the original data with an additional attribute boolean flag indicating whether the link is valid or not, Fil is the filtered data with instances removed for which no link is found, and Att bypasses the original Att input port.

Assume that the original set contains an entity which cannot be resolved in DBpedia, i.e., “MannheimX”. Linking that dataset to DBpedia with the simple linker results in the following table:

ExampleSet (4 examples, 1 special attribute, 4 regular attributes)					
Row No.	id	Code	City	City_link_to_DBpedia	RecordExistsForDBpedia
1	1	67435	Darmstadt	<a href="http://dbpedia.org/resource/Darmstadt">http://dbpedia.org/resource/Darmstadt</a>	true
2	2	68321	Mannheim	<a href="http://dbpedia.org/resource/Mannheim">http://dbpedia.org/resource/Mannheim</a>	true
3	3	62321	MannheimX	<a href="http://dbpedia.org/resource/MannheimX">http://dbpedia.org/resource/MannheimX</a>	false
4	4	165321	Munich	<a href="http://dbpedia.org/resource/Munich">http://dbpedia.org/resource/Munich</a>	true

The output produced by the Web Validator operator then filters the table as follows:

ExampleSet (3 examples, 1 special attribute, 3 regular attributes)				
Row No.	id	Code	City	City_link_to_DBpedia
1	1	67435	Darmstadt	<a href="http://dbpedia.org/resource/Darmstadt">http://dbpedia.org/resource/Darmstadt</a>
2	2	68321	Mannheim	<a href="http://dbpedia.org/resource/Mannheim">http://dbpedia.org/resource/Mannheim</a>
3	4	165321	Munich	<a href="http://dbpedia.org/resource/Munich">http://dbpedia.org/resource/Munich</a>

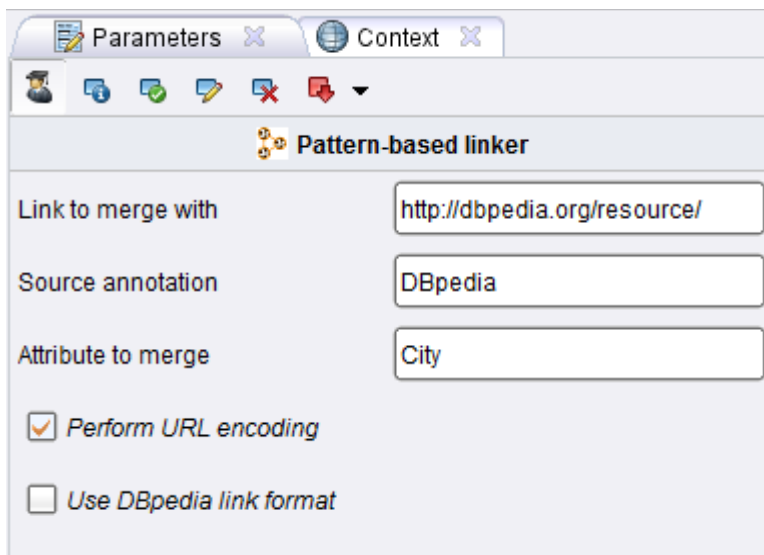
Observation: the entity MannheimX, for which the DBpedia link is wrong, is removed from the dataset.

### 5.1 The Pattern-based Linker

The Pattern-based linker provides a simple method for creating links, which simply guesses links by forming simple URI patterns. The linker takes three arguments:

1. The link prefix
2. The source annotation
3. The attribute to link

Links are created by concatenating the link prefix with the attribute to link. The source annotation field is used for creating a name for the attribute containing the link.



Parameters Context

**Pattern-based linker**

Link to merge with:

Source annotation:

Attribute to merge:

☒ Perform URL encoding

☐ Use DBpedia link format

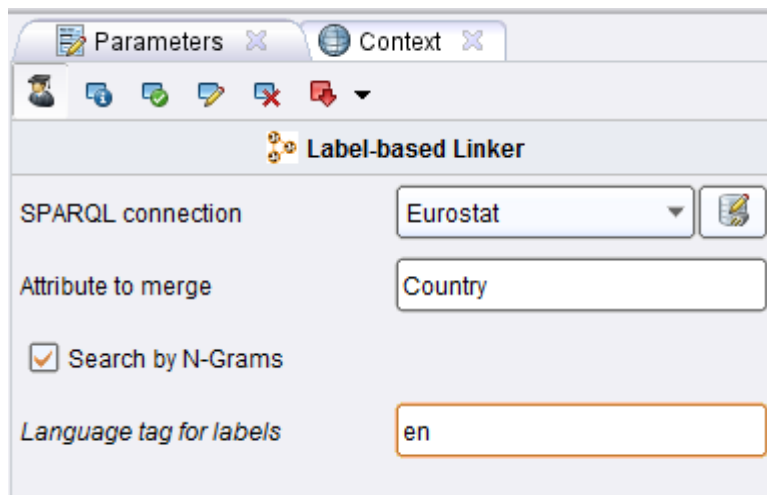
Further attributes that can be set are whether to perform URL Encoding (i.e., replacing special characters and creating proper UTF-8 links), and whether to use a specific link format for DBpedia (e.g., representing blanks as underscores).

In the example shown above, a column “City” containing values such as “Mannheim”, “Darmstadt” etc. would result in the corresponding links “<http://dbpedia.org/resource/Mannheim>”, “<http://dbpedia.org/resource/Darmstadt>”, etc.

## 5.2 The Label-based Linker

The Label-based Linker identifies entities by issuing SPARQL queries against an endpoint, trying to find entities whose labels contain a given string. It takes four parameters:

1. A SPARQL endpoint
2. The attribute to merge (see Pattern-based linker)
3. A flag indicating whether the search should include n-grams.



The screenshot shows a software window titled "Label-based Linker" with two tabs: "Parameters" and "Context". The "Parameters" tab is active. It contains the following fields and controls:

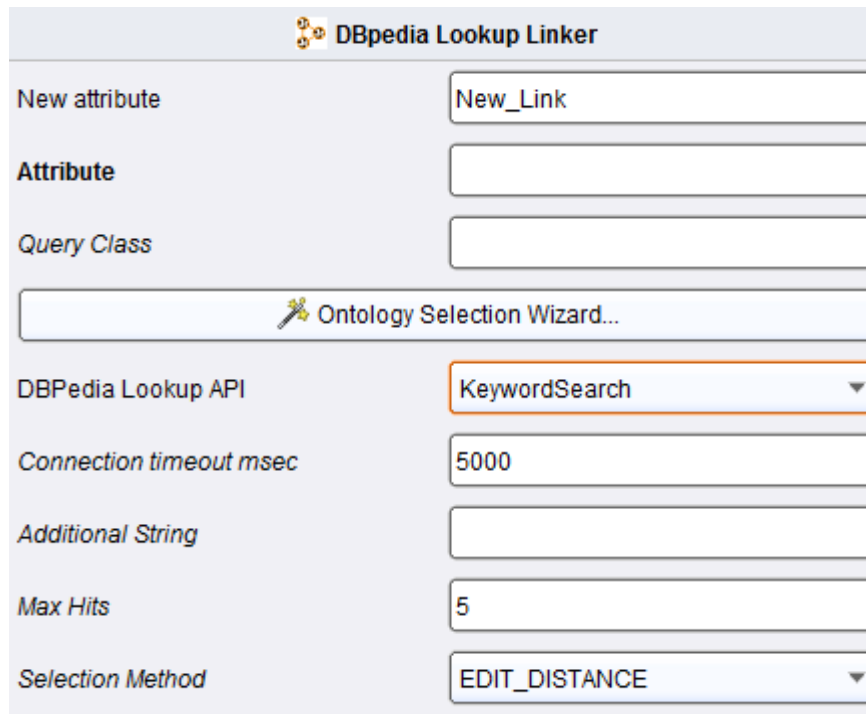
- SPARQL connection:** A dropdown menu showing "Eurostat" with a small icon to its right.
- Attribute to merge:** A text input field containing the word "Country".
- Search by N-Grams:** A checkbox that is checked.
- Language tag for labels:** A text input field containing the code "en".

When including n-grams in the search, the linker will first try to identify entities whose labels contain the string in its entirety, e.g., “United States of America”. If this does not succeed, it will continue with sub tokens, eventually also finding entities whose label is only “United States”.

Note that the SPARQL Based linker, in particular when using the “Search by N-Grams” option, can take some time to link a dataset, since it performs full text search via SPARQL, which, depending on the endpoint used, may not be performed very fast.

In expert mode, you can also restrict the search to labels with a certain language tag, e.g., “en”. If the attribute is left empty (which is the default), labels in all languages, as well as without any language tag, are included in the search.

### 5.3 The DBpedia Lookup Linker



The screenshot shows the 'DBpedia Lookup Linker' configuration window. It has a title bar with a logo and the text 'DBpedia Lookup Linker'. Below the title bar, there are several input fields and a button. The fields are labeled 'New attribute', 'Attribute', 'Query Class', 'DBpedia Lookup API', 'Connection timeout msec', 'Additional String', 'Max Hits', and 'Selection Method'. The 'New attribute' field contains the text 'New\_Link'. The 'Attribute' field is empty. The 'Query Class' field is empty. Below these fields is a button labeled 'Ontology Selection Wizard...'. The 'DBpedia Lookup API' field is a dropdown menu with 'KeywordSearch' selected. The 'Connection timeout msec' field contains the value '5000'. The 'Additional String' field is empty. The 'Max Hits' field contains the value '5'. The 'Selection Method' field is a dropdown menu with 'EDIT\_DISTANCE' selected.

DBpedia Lookup<sup>5</sup> is a web service which allows search over DBpedia. It can be used as an operator for linking and is more versatile than the pattern-based and the label-based linker.


There are various parameters for the lookup linker:

- *New Attribute* is the name of the attribute containing the link
- *Attribute* is the name of the attribute that contains the text value
- *QueryClass* is an optional parameter which can be used to narrow the search results. It is a class name from the DBpedia ontology<sup>6</sup>, e.g., “City”. If the attribute to link contains city names, and *QueryClass* is not set, the linker will link the cities, e.g., <http://dbpedia.org/resource/Mannheim>. If the *QueryClass* is set to “University” instead, the link will be [http://dbpedia.org/resource/University\\_of\\_Mannheim](http://dbpedia.org/resource/University_of_Mannheim) instead. For convenience, the class can be selected from a visualized type hierarchy of the DBpedia ontology:

---

<sup>5</sup><http://lookup.dbpedia.org/>

<sup>6</sup><http://dbpedia.org/Ontology>



- *DBpedia Lookup API* defines the URL of the lookup service. There are two predefined endpoints (the keyword and prefix lookup<sup>7</sup>, as well as the possibility to define a user-specified service URL)
- *Connection timeout msec* defines the timeout interval of the service.
- *Additional string* can be used to define any postfix that is respected when using the EDIT\_DISTANCE selection method (see below). For example, if “university” is added, results containing “university” in their name are preferred.
- *Max Hits* defines the maximum number of results used from spotlight (per query).
- *Selection Method* defines how to disambiguate between multiple results. EDIT\_DISTANCE, JARO\_WINKLER and JACCARD select the result whose name is closest to the original attribute using the respective string similarity function<sup>8</sup>. FIRST\_RESULT uses the first result, which is usually the most prominent one.

---

<sup>7</sup>See <http://lookup.dbpedia.org/> for details.

<sup>8</sup>See [http://en.wikipedia.org/wiki/String\\_metric](http://en.wikipedia.org/wiki/String_metric) for details.

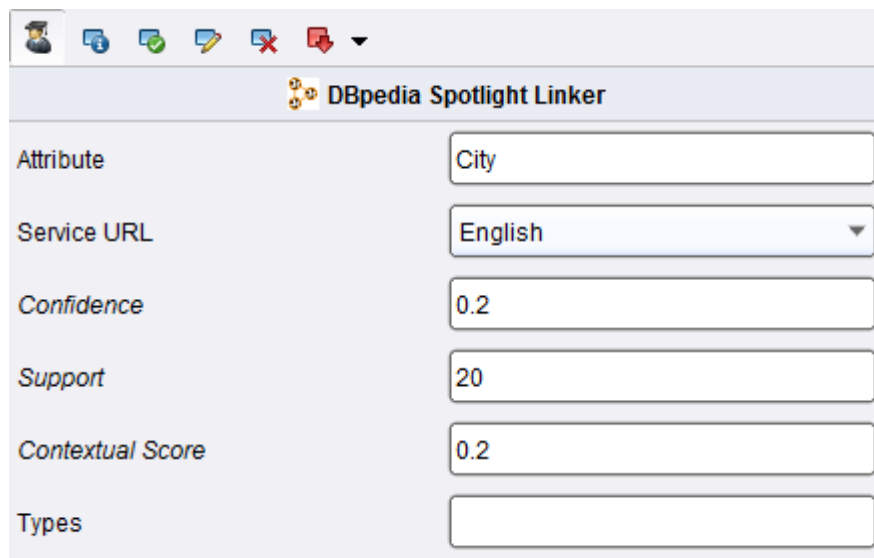
## 5.4 The DBpedia Spotlight Linker

The DBpedia spotlight linker is a linker which can be used for linking text. It produces a set of attributes, each containing one concept identified. For example, for an attribute “text”, the attributes “text\_Concept\_1”, “text\_Concept\_2”, etc., are created.

In subsequent generators, the features generated for those attributes are automatically merged. For example, with a simple type feature generator, there would only be one feature called text\_Concept\_1\_City, no matter if the city is identified as the first, second, third, etc. concept in the text.

The spotlight linker allows for configuring parameters of the service – confidence, support, and contextual score. Details can be found in the documentation of DBpedia Spotlight.<sup>9</sup>

The *Service URL* parameter allows for selecting the URL of the service. Some public installed services for common languages are predefined, “Custom” allows for specifying another URL.



DBpedia Spotlight Linker	
Attribute	City
Service URL	English
Confidence	0.2
Support	20
Contextual Score	0.2
Types	

The *Types* parameter allows for finding only a list of types (e.g., Persons, Companies, etc.), as specified in the DBpedia ontology (see above) as a comma separated list.

## 5.5 The SameAs Linker

In Linked Open Data, a common way to establish links between datasets is the use of owl:sameAs links, indicating that a resource in dataset 1 is equal to a resource in dataset 2. The owl:sameAs linker follows such links in order to establish links to other data sources.

The typical use case for using the SameAs linker is to create links for a dataset that is difficult (or too slow) to link to with the other approaches, but which provides links to DBpedia, where links can be created, e.g., with the DBpedia Lookup Linker. Consider a setting where the Eurostat Linked Open Data endpoint contains triples such as


```
eurostat:France owl:sameAs dbpedia:France .
```

These links can be used using the following setup:

---

<sup>9</sup><http://spotlight.dbpedia.org/>





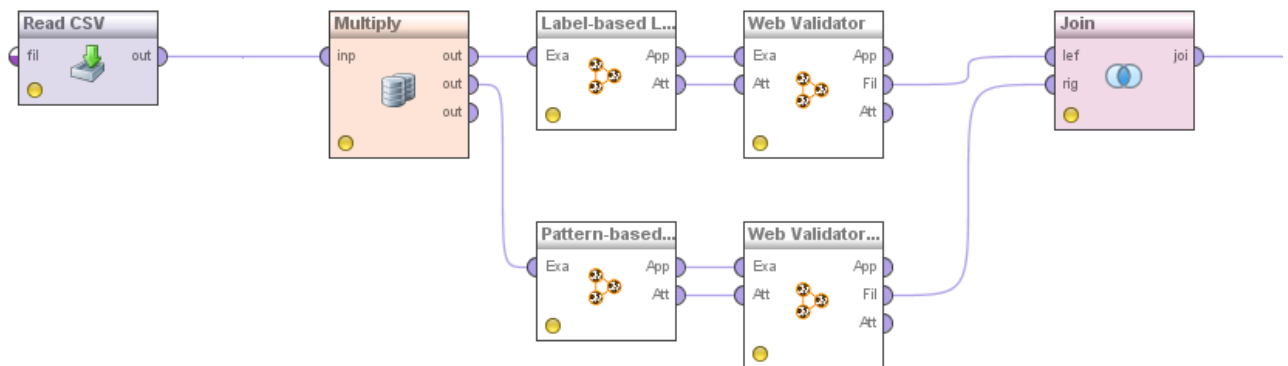
For configuring the linker, the user has to specify the field containing the original link, the name for the new link, the SPARQL endpoint used to search for owl:sameAs links, and optionally a string pattern for filtering the results:

The screenshot shows the 'SameAs Linker' configuration window. It has a title bar with the component icon and name. Inside, there are four fields: 'SPARQL connection' with a dropdown menu set to 'Eurostat' and a refresh icon; 'Attribute to search for' with a text box containing 'DBpedia'; 'New attribute name' with a text box containing 'Eurostat'; and 'Endpoint search pattern' with an empty text box.

With the pipeline and setting shown above, Eurostat URIs are found by first searching for DBpedia URIs, and then exploiting sameAs links in Eurostat.

## 5.6 Combining Linkers

In some cases, it might be useful to combine linkers – e.g., to add information from multiple sources, or to add information about different entities in a dataset, e.g., both the film as well as its director. In that case, the data can be multiplied with the Multiply operator (in “Process control”), and the output of both linkers is joined with the Join operator (in “Data Transformation”). Note that for the latter step, the original input data needs to define at least one ID column.



## 6 Using generators

The RapidMiner Linked Open Data extension supports different generators, which perform different knowledge extraction and aggregation steps.

A generator needs to know which attribute(s) contain the link(s). There are two ways of passing that information to the operator:

- by wiring the Att output of a linker to the corresponding input of the generator (this is the standard way when using generators).
- by entering the attribute name(s) in the Attribute to extend parameter (this is the way to go if the dataset at hand already contains links).

Apart from the attribute generation, most generators also generate a hierarchy of the attributes generated, which can be used for later filtering (see below).

Furthermore, each generator requires the selection of a SPARQL endpoint. For using access based on URI dereferencing, select a default SPARQL endpoint configured with the option “Use URL data”.

## 6.1 Direct Types Generator

The Direct Types Generator extracts all statements with predicate `rdf:type` and creates a boolean attribute for each type. For example, from a statement like

```
<http://dbpedia.org/resource/Mannheim> rdf:type  
<http://dbpedia.org/ontology/City> .
```

an attribute with name City\_link\_to\_DBpedia would be created, and the value for the instance linked to <http://dbpedia.org/resource/Mannheim> would be set to true.

The process using the Simple Types Generator looks as follows:



Note that both outputs of the linker – App and Att – need to be wired to the respective input ports of the generator. The output looks as follows:

[illegible]

## 6.2 Data Properties Generator

The Data Properties Generator creates an attribute for each literal value that the linked instances has. For example, from the property

```
<http://dbpedia.org/resource/Mannheim>  
<http://dbpedia.org/ontology/population> "123873"^^xsd:integer .
```

the generator would create an attribute

City\_link\_to\_DBpedia\_data\_http://dbpedia.org/ontology/population and set the value for the instance linked to Mannheim to 123873. Some basic guessing of data types is performed, e.g., numeric values are marked as such, and numeric literals without an explicit type are parsed into numbers, if possible.

## 6.3 Unqualified and Qualified Relation Generators

The RapidMiner Linked Open Data extension supports generators exploiting relations to other resources. The *unqualified* relations generators create attributes from the existence of relations, while the *qualified* relations generators also take the types of the related resources into account. Both generators exist in two variants, one generating boolean, one generating numeric features.

Consider the following triples:

```
<http://dbpedia.org/resource/Berlin>  
<http://dbpedia.org/ontology/capitalOf>  
<http://dbpedia.org/resource/Germany> .  
  
<http://dbpedia.org/resource/Germany> rdf:type  
<http://dbpedia.org/ontology/Country> .
```

The Relation Boolean generator would create an attribute

City\_link\_to\_DBpedia\_in\_boolean\_http://dbpedia.org/ontology/capitalOf and set its value to *true* for an instance linked to <<http://dbpedia.org/resource/Berlin>> . The Relation Numeric generator would create an attribute City\_link\_to\_DBpedia\_in\_numeric\_http://dbpedia.org/ontology/capitalOf and set its value to *1* for an instance linked to <<http://dbpedia.org/resource/Berlin>>.

The Qualified Relation Boolean generator would create an attribute

City\_link\_to\_DBpedia\_in\_boolean\_http://dbpedia.org/ontology/capitalOf\_type\_http://dbpedia.org/ontology/Country and set its value to *true* for an instance linked to <<http://dbpedia.org/resource/Berlin>> . The

Qualified Relation Numeric generator would create an attribute

City\_link\_to\_DBpedia\_in\_numeric\_http://dbpedia.org/ontology/capitalOf\_type\_http://dbpedia.org/ontology/Country and set its value to *1* for an instance linked to <<http://dbpedia.org/resource/Berlin>>.

Note that in especially the qualified relation generators, while very powerful, may create a very large number of attributes, which can cause problems for subsequent processing steps.

## 6.4 The Specific Relation Generator

In order to pursue only a specific predicate, the specific relation linker can be used. Optionally, a predicate to use for generating the hierarchy can be specified.

For example, the configuration below extracts all categories from DBpedia, and uses the skos:broader predicate (which is used for category hierarchies in DBpedia) for generating the hierarchy:

**Specific Relation Generator**

SPARQL connection: DBpedia

☒ Generate Hierarchy

Direct Relation: <http://purl.org/dc/terms/subject>

Hierarchy Relation: <p://www.w3.org/2004/02/skos/core#broader>

Max hierarchy depth: 1

## 6.5 The Custom SPARQL Generator

In cases where you know the dataset you are using and want to add specific data not covered by any of the default generators, you can also define your own SPARQL statements. For example, for a dataset of cities, you might want to add the population of the state that the city is in. This can be done with the custom SPARQL generator, which allows to create a custom statement:

**Edit Parameter Text: SPARQL query**

Edit Parameter Text: SPARQL query  
The SPARQL Query.

```

1 SELECT ?populationState WHERE {
2   *city_link_to_DBpedia* <http://dbpedia.org/ontology/federalState> ?state .
3   ?state <http://dbpedia.org/ontology/populationTotal> ?populationState .
4 }

```

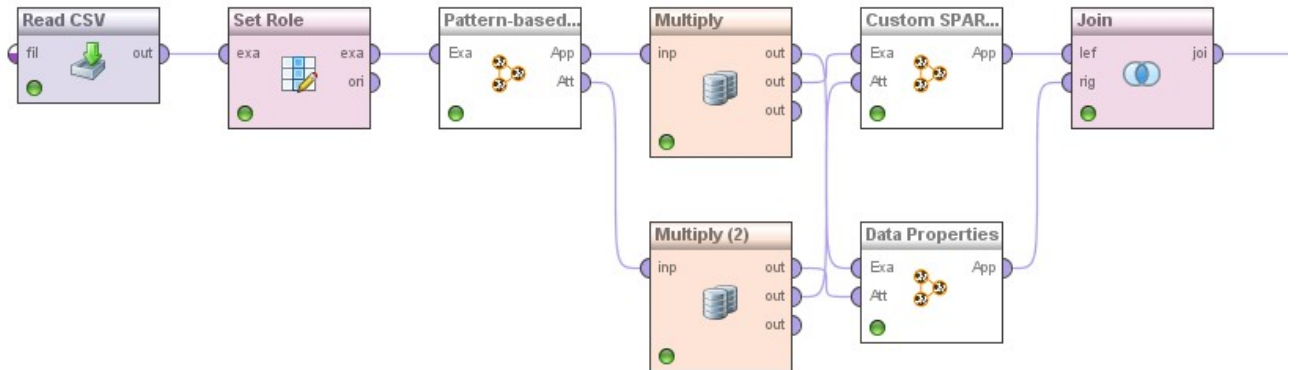
Enlarge Apply Cancel

Within the statement, you can use links generated by linkers as placeholders, enclosed in asterisks. In the above example, a linker has generated the link attribute `city_link_to_DBpedia`, which is used in the query. The generator then generates a new attribute called “populationState” (using the variable name for naming the attribute) and adds it to the dataset.

The reserved term `*ATT_INPUT_PORT*` can be used in the query. It is filled with the attributes passed at the att input port.

## 6.6 Combining Generators

For combining generators on the same link, e.g., extracting both data properties and types, the following process setup has to be used:



Two Multiply operators are used, one for multiplying the data, one for multiplying the attribute information. Like for combining linkers, the input data needs to define at least one ID attribute to allow the final Join. The Set Role operator can be used for defining IDs.

Note that you can also combine generators with data from different data sources. To that end, you have to run a linker and one or more generators for each data source.

## 7 Filtering Attributes


The RapidMiner LOD extension may create a lot of attributes. Thus, it is advisable to filter the attributes. Besides standard attribute selection operators that exist in RapidMiner (e.g. Remove Correlated Attributes), the extension provides means to exploit the semantics of Linked Open Data sources used to filter attributes.

Some generators, such as the direct types generator, also create a hierarchy of attributes. As for direct types, `rdfs:subClassOf` links are exploited to build the hierarchies of attributes generated.

The Simple Hierarchy Filter operator uses that hierarchy to filter attributes. Two modes are available:

- *Filter by correlation.* If two attributes A and B are in a hierarchy relation  $A < B$ , and the correlation of A and B is larger than a threshold, then A is removed, and only B is kept.
- *Filter by information gain.* If two attributes A and B are in a hierarchy relation  $A < B$ , and the difference in information gain between A and B is smaller than a threshold, then A is removed, and only B is kept. Note that this variant requires a (nominal) label in the dataset.

Both modes avoid creating too specific features, and learning overfitting models. The usage is shown below.



## 8 Storing Results

For larger datasets, the runtimes of linkers and generators can be quite high, depending on the dataset and specific operators used. Since in a typical workflow, you will first generate data and then experiment with a larger number of subsequent operators (e.g., classification or regression algorithms), we suggest that you store the output dataset of the LOD operator chain in an intermediate file (e.g., CSV) or database, and perform all subsequent analysis on that stored dataset. That way, you will have to run the time consuming access to Linked Open Data only once.

## 9 Migration from Earlier Versions of the Extension

Migrating from earlier versions of the extension may cause problems in the SPARQL configurations dialog. If you experience such problems (which manifest themselves by repeated questions whether you want to store a connection), locate and delete the file `configurable-sparqlconfig.xml` on your harddisk, and restart RapidMiner.



## 10 Release Notes

Version 1.3.3, released on May 15<sup>th</sup>, 2014

- Additional configuration options for DBpedia Lookup and DBpedia Spotlight
- Performance optimizations
- Bug fixes

Version 1.3.2, released on March 14<sup>th</sup>, 2014

- The SPARQL importer and generator now support common prefixes (e.g., rdf, owl, foaf, etc.) through prefix.cc
- Bugfix in SPARQL importer

Version 1.3.1, released on March 7<sup>th</sup>, 2014

- Support for feature extraction by resolving URIs
- Entity class detection during linking
- Visual selection of the query class for the DBpedia Lookup linker
- Performance improvements through multi threading
- Bugfixes

Version 1.2.116, released on November 29<sup>th</sup>, 2013

- Specific relation generator
- Sameas linker
- Simple hierarchy filter
- Support for pre-defined links
- Performance improvements
- Bugfixes

Version 1.1.024, released on October 15<sup>th</sup>, 2013

- Support for local RDF files, including reasoning
- DBpedia Lookup linker
- DBpedia Spotlight linker
- Support for binary features in Custom SPARQL generator
- Bugfixes

Version 1.0.071, released on September 13<sup>th</sup>, 2013.